

I. Pen-and-paper

1) [4v] Draw the training confusion matrix.

		Previstos	
		P	N
Reais	P	8	3
	N	4	5

$$\#P_{reais} = 11, \quad \#N_{reais} = 9,$$


$$\#P_{previstos} = 12, \quad \#N_{previstos} = 8, \quad \#total = 20.$$

2) [3v] Identify the training F1 after a post-pruning of the given tree under a maximum depth of 1.

$$- F1 = Fmeasure_1 = F_{\beta=1} = F_{\alpha=0.5} = \frac{1}{\frac{1}{2} \frac{1}{P} + \frac{1}{2} \frac{1}{R}} = \frac{2}{\frac{1}{P} + \frac{1}{R}};$$

$$accuracy = \frac{\#bem\ classificados}{\#total\ classificados}, \quad recall/sensitivity = \frac{TP}{TP + FN}, \quad precision = \frac{TP}{TP + FP}.$$

Decision tree original:




$$Precision_P = \frac{8}{8+4} = \frac{8}{12},$$

$$Recall_P = \frac{8}{8+3} = \frac{8}{11}, \quad F1_P = \frac{2}{\frac{12}{8} + \frac{11}{8}} = \frac{16}{23}.$$

$$Precision_N = \frac{5}{5+3} = \frac{5}{8},$$

$$Recall_N = \frac{5}{4+5} = \frac{5}{9}, \quad F1_N = \frac{2}{\frac{8}{5} + \frac{9}{5}} = \frac{10}{17}.$$

Decision tree depois do post-pruning:



		Previstos	
		P	N
Reais	P	5	6
	N	2	7

- confusion matrix da nova decision tree

$$Precision_P = \frac{5}{5+2} = \frac{5}{7},$$

$$Recall_P = \frac{5}{5+6} = \frac{5}{11}, \quad F1_P = \frac{2}{\frac{7}{5} + \frac{11}{5}} = \frac{5}{9}.$$

$$Precision_N = \frac{7}{7+6} = \frac{7}{13},$$

$$Recall_N = \frac{7}{7+2} = \frac{7}{9}, \quad F1_N = \frac{2}{\frac{13}{7} + \frac{9}{7}} = \frac{7}{11};$$

3) [2v] Identify two different reasons as to why the left tree path was not further decomposed.

O caminho esquerdo da árvore não foi mais desenvolvido para evitar overfitting de duas formas (pelo menos):

1. Sendo o data split não estatisticamente significativo, ao pararmos a expansão do nó, evitamos a produção de filhos que se baseiam em samples muito pequenos;
2. A expansão do nó não respeitaria o valor mínimo de diminuição do seu nível de impureza (definida no parâmetro *min_impurity_decrease* na função *sklearn.tree.DecisionTreeClassifier*, com default = 0).

4) [3v] Compute the information gain of variable y1.

$$IG(out | y1) = H(out) - H(out | y1);$$

$$H(out) = - \sum_{v \in out} P(v) \log_2(P(v));$$

$$H(out | y1) = - \sum_{v \in y1} P(y1 = v) H(out | y1 = v), \quad \text{sendo out a variável de output.}$$

$$- H(out) = - \frac{11}{20} \log_2\left(\frac{11}{20}\right) - \frac{9}{20} \log_2\left(\frac{9}{20}\right) \approx 0.99277$$

$$- H(out | y1) = \frac{7}{20} \left[- \frac{5}{7} \log_2\left(\frac{5}{7}\right) - \frac{2}{7} \log_2\left(\frac{2}{7}\right) \right] + \frac{13}{20} \left[- \frac{6}{13} \log_2\left(\frac{6}{13}\right) - \frac{7}{13} \log_2\left(\frac{7}{13}\right) \right] \approx 0.94932$$

$$\text{Logo, } IG(out | y1) = H(out) - H(out | y1) = 0.99277 - 0.94932 = 0.04345 \quad \square$$

II. Programming and critical analysis

5) Código python:

```
# Import wall

import pandas as pd

import numpy as np

from sklearn import metrics, datasets, tree

from scipy.io.arff import loadarff

from sklearn.model_selection import train_test_split

from sklearn.feature_selection import mutual_info_classif, SelectKBest

import matplotlib.pyplot as plt


# Reads the 'pd_speech.arff' file and creates the desired data frame.

data = loadarff('pd_speech.arff')

df = pd.DataFrame(data[0])


# Discretization of the output variable 'class' (B1 -> 1).

df['class'] = df['class'].str.decode('utf-8')


# Defines the X and Y data sets:

x = df.drop("class", axis=1)

y = np.ravel(df['class'])


num_features = [5, 10, 40, 100, 250, 700]

acc_train = []

acc_test = []


for n in num_features:

    # Feature selection:

    # - selects the k best variables based on the

    #   mutual info classifier (information gain).

    selector = SelectKBest(mutual_info_classif, k=n)
```

Aprendizagem 2022/23
Homework I – Group 018

Continuação (II. 5):

```
# Gets the sub-set to be used on this iteration,  
# based on the selector created above.  
x_reduced = selector.fit_transform(x, y)  
  
# 70-30 training-testing split.  
x_train, x_test, y_train, y_test = train_test_split(x_reduced, y, test_size = 0.3,  
                                                    random_state=1)  
  
# Trains the decision tree on the training sets.  
decision_tree = tree.DecisionTreeClassifier()  
predictor = decision_tree.fit(x_train, y_train)  
  
# Gets the predicted values for both the train  
# and test samples.  
y_train_pred = predictor.predict(x_train)  
y_test_pred = predictor.predict(x_test)  
  
# Test model performance:  
# - calculates the accuracy of both samples,  
#   train and test.  
acc_train += [round(metrics.accuracy_score(y_train, y_train_pred), 2)]  
acc_test += [round(metrics.accuracy_score(y_test, y_test_pred), 2)]  
  
print(acc_train)  
print(acc_test)
```

output:

```
[1.0, 1.0, 1.0, 1.0, 1.0, 1.0]  
[0.75, 0.75, 0.84, 0.81, 0.82, 0.79]
```

Aprendizagem 2022/23
Homework I – Group 018

6) [2v] Why training accuracy is persistently 1? Critically analyze the gathered results.

Pela sua definição, accuracy corresponde ao número de previsões corretas a dividir pelo número total de previsões.

No caso concreto da accuracy de treino, todas as previsões serão corretas (dando sempre $\text{accuracy} = 1$), uma vez que são utilizadas tanto para treinar como para testar o próprio modelo usado no training set (70%).

Da análise dos resultados obtidos, podemos concluir que os modelos criados se adequam ao nosso dataset, com uma accuracy média de 80% (com um training-testing split 70-30).

